

Sebastian Sardina <[ssardina@gmail.com](mailto:ssardina@gmail.com)>

---

**Good news: IJCAI 2011 submission number 477**

---

Toby.Walsh@nicta.com.au <[Toby.Walsh@nicta.com.au](mailto:Toby.Walsh@nicta.com.au)>

Thu, Mar 31, 2011 at 11:13 PM

To: [sebastian.sardina@rmit.edu.au](mailto:sebastian.sardina@rmit.edu.au)

Dear Dharendra Singh:

I am delighted to inform you that the following submission has been accepted to appear in the IJCAI 2011 proceedings, and to be presented at the conference by means of both a talk and a poster:

Integrating Learning into a BDI Agent for Environments  
with Changing Dynamics

The Program Committee did an amazing job reviewing the 1325 papers submitted to the conference. Of these, only 227 papers (17%) were considered to be of sufficient quality and interest to merit both a talk and a poster. Please repay their efforts by taking on board their suggestions when you revise your paper. You have two weeks to prepare your final copy.

The deadline for uploading camera ready copy is 15 April 2011 at <http://www.aaai.org/Publications/Author/registration-letter-ijcai.php>  
I will send you more detailed instructions shortly (including how you can if you wish purchase two additional pages to the six you already have).

There are several other deadlines in the next few weeks that may be relevant to you.

First, many workshops and related events have their submission deadlines in the next week or two. Please consider participating in these. See <http://ijcai-11.iiia.csic.es/program/workshops>

IJCAI and the AI Journal offer travel grants to attend IJCAI. The deadline for application is April 7th: [http://ijcai-11.iiia.csic.es/travel\\_grants](http://ijcai-11.iiia.csic.es/travel_grants)

There is also a volunteer program (deadline April 15th) which provides complementary registrations: [http://ijcai-11.iiia.csic.es/students/volunteers\\_program](http://ijcai-11.iiia.csic.es/students/volunteers_program)

Registration for IJCAI 2011 will go live in the next week. Early registration is by May 15th.

Finally, let me congratulate you on your fine work. If you have any additional questions, please feel free to get in touch.

Best Regards,  
Toby Walsh  
IJCAI 2011

## =====

IJCAI 2011 Reviews for Submission #477

=====

Title: Integrating Learning into a BDI Agent for Environments with Changing Dynamics

Authors: Dharendra Singh, Sebastian Sardina, Lin Padgham and Geoff James

=====

## REVIEWER #1

=====

-----  
Reviewer's Scores

-----

Relevance: 10  
Significance: 9  
Soundness: 9  
Originality: 9  
Evaluation and comparison: 10  
Clarity: 10  
OVERALL SCORE: 9

-----  
Comments

-----

## Relevance:

The paper is extremely relevant to the special track on Integrated and Embedded Artificial Intelligence. It presents an integration of machine learning techniques into a BDI Agent Model, which uses a measure of confidence to determine whether the learned plan trees for behaviour are appropriate (favoring exploitation), or not (favoring exploration of new strategies).

## Significance and Originality:

One of the interesting aspects of the work presented is that the confidence measure does not converge monotonically to 1, making the agent always to perform exploitation. Departing from traditional approaches, authors assume that the learned behavior can sometimes become wrong and inappropriate in a dynamic environment where rules are not constant. To solve this, authors put forward a dynamic confidence measure that adapts accordingly with changes in performance.

## Soundness:

The work presented, to the best of my knowledge seems to be sound and correct. I managed to follow and understand the formulas and formal definitions presented. However, since my background is not in machine learning, it is hard for me to properly evaluate their soundness.

## Evaluation and comparison:

The evaluation performed seems very appropriate and I think it managed to prove

the proposed plan-selection learning mechanism is able to adapt when the dynamics of the environment change over time.

Clarity :

The paper is clear and well written, with a clear description of the research goal and their contribution to the area. Although formal in some parts, authors manage to make the paper quite easy to read and follow by providing a set of illustrative and meaningful examples.

=====

REVIEWER #2

=====

-----

Reviewer's Scores

-----

Relevance: 9  
Significance: 8  
Soundness: 8  
Originality: 8  
Evaluation and comparison: 6  
Clarity: 8  
OVERALL SCORE: 7

-----

Comments

-----

A BDI plan-selection learning mechanism is described that uses a dynamic confidence measure to judge an agent's current understanding (of how well plans solve goals) and the need for exploration to improve that understanding. The mechanism is designed for dynamic environments where the confidence in a plan does not increase monotonically. The work builds on the BDI learning approach of Singh, et al., and preserves the monotonic convergence of that approach in static environments.

I liked this paper. The extensions to BDI plan-selection learning are interesting and important, and the paper is generally clear and well written. The experiments section presents the behavior of using the extended mechanism in a modular battery controller setting. The presented results show the capabilities of the mechanism, but I would have liked to see a comparison with the performance obtained (or not) using a Singh-like monotonic learning mechanism in the same dynamic setting.

The abstract needs improvement/specificity (such as including some of the summary information in the first paragraph above as well as the contributions claimed).

Make clear that the "they" in the second sentence of the Introduction refers to the "classic" or "original" BDI paradigm (as when I first read the paper, I didn't realize until later in the Introduction that the authors were even aware of other BDI learning work).

I wanted more discussion of value(s) for "n" in the confidence measure (and

indirectly in the final plan selection weight). For example, what "n" value(s) was used in your experiments? What adjusts "n" if the rate of dynamic change changes (what is the effect of using an inappropriate value of "n")? Was any thought given to using temporally reduced weighting (such that old experience carries less weight) rather than the "n" based cutoff with uniform weighting?

Minor nit: eliminate the space between "trace" and footnote reference 1.

The citations and discussion of other work are reasonable, as is the mention of limitations and possible work in the Conclusion. Again, I liked this paper.

=====

REVIEWER #3

=====

-----

Reviewer's Scores

-----

Relevance: 7  
Significance: 4  
Soundness: 6  
Originality: 4  
Evaluation and comparison: 4  
Clarity: 7  
OVERALL SCORE: 4

-----

Comments

-----

SUMMARY

The work in this paper extends the existing BDI learning framework: it proposes a heuristic (the equation towards the end of Section 3.2) measuring the confidence of plans that automatically adjusts with the environment dynamics. The running examples are great for our understanding of the authors' work. I'm especially concerned by the significance and thoroughness of the experimental results presented in this paper. This work can be improved based on my feedback below:

DETAILED REVIEW

ORIGINALITY:

The authors propose a minor fix to the existing BDI learning framework so that it can be used in a dynamically changing environment. The proposed simple idea does contribute incrementally towards agent programming research (but not towards machine learning research, as indicated by the authors). See the next few comments to see how this work can be improved.

SIGNIFICANCE, EVALUATION, AND COMPARISON:

(1) The authors have mentioned the importance of addressing the exploration vs. exploitation dilemma in section 2.2 using the confidence measure. However, I notice that only the exploitation part is addressed in section 3.3. The exploration part seems to be missing, which, if included, will raise the significance of this work.

(2) I'm particularly concerned by the thoroughness of the experimental results, which can be significantly improved:

(a) From Fig. 2, it appears that the learning time is long. I recommend that the authors conduct experiments or provide theoretical results analyzing the tradeoff(s) between the degree of change of environment dynamics and speed of learning and size of state space.

(b) The authors did not provide experimental results comparing its proposed technique with the state-of-the-art BDI learning techniques. I recommend that the authors compare with the one-off learning techniques of Singh et al. (2010a), (2010b) by introducing a fixed re-learning rate.

#### TECHNICAL SOUNDNESS:

Thanks for clarifying the doubts that I have, as stated below.

(1) I understand from the paragraphs before Algorithm 1 that the stability degree is recorded in a plan only for failed execution traces. If this is the case, is it possible for the stability degree  $\lambda$  to converge to 1 for fixed dynamics (as indicated in the paragraph after Algorithm 1) without taking into account the successful execution traces and the successful stable (possibly non-leaf) plans?

(2) The authors have said that "When a plan finally succeeds, we take an optimistic view and record degree 1 against it". Does this apply to non-leaf plans that succeed (though they may fail in other execution traces) as well?

(3) In footnote 3, I'm surprised to see that the value of  $\alpha$  is not critical to performance. If I set  $\alpha$  to either of the extremes (0 or 1), one of the confidence measures is dispensable.

(4) It would be clearer if the authors elaborate on the relationship between the RecordDegreeStability function in Algorithm 1 and the average degree of stability in the text. Otherwise, a reader may be mistaken from Algorithm 1 that the current stability degree is entirely replaced by the newly-computed stability degree for plan P1 instead of computing their weighted average.

(5) The authors say that "NewStates(P, n) is the set of world states in the last n executions of P that have also been witnessed before that". Shouldn't this be corrected to "... have NOT been witnessed..."?

#### CLARITY:

(1) This paper is clearly written except for some parts in Section 3.2.

(2) The last paragraph of the conclusion should be put in front as a separate section on related work.

---